

Haekyu Park

haekyu@gatech.edu · haekyu.com · Curriculum Vitae · Google Scholar

My research goal is to enhance **machine learning interpretability**, to enable:

- **Confident use of machine learning**, by helping people interpret and trust the models
- **Effective model debugging and improvement**, by promoting practitioners' understanding of why a model would not work well and providing insights into how to fix it
- **Broadening access to machine learning**, by encouraging people of diverse backgrounds to easily learn and leverage the technologies.

Specifically, I design and develop interactive visualization systems to interpret machine learning models. I am a member of the [Polo Club of Data Science](#) at Georgia Tech. I have been fortunate to work with amazing researchers, engineers, and scientists at NVIDIA and Intel.

Education

Georgia Institute of Technology

Ph.D., Computer Science

Advisor: Dr. [Polo Chau](#)

Aug 2018 - Present

Seoul National University

B.S., Computer Science and Engineering

Graduated with honors (Cum Laude)

Mar 2012 - Aug 2017

Research Experience

AI Infrastructure Software Intern

May 2020 - Jul 2020

NVIDIA, Santa Clara, CA

Mentor: Joe Eaton, Brad Rees, Bartley Richardson

Developed a visual graph analytics, allowing for interactively running multiple graph algorithms in real-time on large graphs.

Leveraged GPU acceleration for both data analysis and rendering side.

Data Science Intern

May 2019 - Aug 2019

NVIDIA, Austin, TX

Mentor: Bartley Richardson, Brad Rees, Joe Eaton

Internship results are integrated into NVIDIA RAPIDS team's [cybersecurity usecase notebook](#), presented at [KDD 2019 NVIDIA RAPIDS tutorial](#).

Graduate Research Assistant	Aug 2018 - Present
Georgia Institute of Technology, Atlanta, GA	
Undergraduate Research Assistant	June 2016 - Aug 2017
Seoul National University, Seoul, Republic of Korea	

Awards & Honors

"Thank a Teacher" Award	2019
Center of Teaching & Learning (CTL), Georgia Institute of Technology	
Moon-Jung Chung Scholarship	2019
KOCSEA (The Korean Computer Scientists and Engineers Association in America)	
National Scholarship for Science and Engineering	2015
National Scholarship for Science and Engineering	

Grants & Funding

WiML Travel Funding	2019
\$550 Travel Funding	
Women in Machine Learning Workshop, co-located with NeurIPS	
Amazon AWS Research Grant	2018
Funded \$5,000 in AWS cloud credits	
Co-PIs: Nilaksh Das, Scott Freitas, Duen Horng Chau	

Publication

SkeletonVis: Interactive Visualization for Understanding Adversarial Attacks on Human Action Recognition Models

[Haekyu Park](#), Zijie J. Wang, Nilaksh Das, Anindya S. Paul, Pruthvi Perumalla, Zhiyan Zhou, and Duen Horng Chau
 AAAI, Demo, Virtual, 2021.
[▶ Demo](#) [PDF](#) [arXiv](#)

Bluff: Interactively Deciphering Adversarial Attacks on Deep Neural Networks

Nilaksh Das*, [Haekyu Park*](#), Zijie J. Wang, Fred Hohman, Robert Firstman, Emily Rogers, and Duen Horng Chau
 IEEE Visualization Conference, ([VIS](#)), Salt Lake City, UT, USA, 2020.
 * Authors contributed equally.
[▶ Demo](#) [PDF](#) [arXiv](#)

CNN Explainer: Learning Convolutional Neural Networks with Interactive Visualization

Zijie J. Wang, Robert Turko, Omar Shaikh, [Haekyu Park](#), Nilaksh Das, Fred Hohman, Minsuk Kahng, Duen Horng Chau
 IEEE Conference on Visual Analytics Science and Technology, ([VAST](#)), Salt Lake City, UT, USA, 2020.
[▶ Demo](#) [PDF](#)

A Comparative Analysis of Industry Human-AI Interaction Guidelines

Austin P. Wright, Zijie J. Wang, [Haekyu Park](#), Grace Guo, Fabian Sperrle, Mennatallah El-Assady, Alex Endert Daniel Keim Duen Horng Chau
IEEE Visualization Conference, Workshop on Trust and Expertise in Visual Analytics (TREX), Salt Lake City, UT, USA, 2020.

[PDF](#)

Argo Lite: Open-Source Interactive Graph Exploration and Visualization in Browsers

Siwei Li, Zhiyan Zhou, Anish Upadhyay, Omar Shaikh, Scott Freitas, [Haekyu Park](#), Zijie J. Wang, Susanta Routray, Matthew Hull, and Duen Horng Chau
ACM International Conference on Information and Knowledge Management, ([CIKM](#)), Resource Track, Online, 2020.

[▶ Demo](#) [PDF](#)

Massif: Interactive Interpretation of Adversarial Attacks on Deep Learning

Nilaksh Das*, [Haekyu Park](#)*, Zijie J. Wang, Fred Hohman, Robert Firstman, Emily Rogers, and Duen Horng Chau
ACM CHI Conference on Human Factors in Computing Systems (CHI), [Late-Breaking Works](#), Honolulu, Hawaii, USA, 2020.
* Authors contributed equally.

[PDF](#) [arXiv](#)

CNN 101: Interactive Visual Learning for Convolutional Neural Networks

Zijie J. Wang, Robert Turko, Omar Shaikh, [Haekyu Park](#), Nilaksh Das, Fred Hohman, Minsuk Kahng, and Duen Horng Chau
ACM CHI Conference on Human Factors in Computing Systems (CHI), [Late-Breaking Works](#), Honolulu, Hawaii, USA, 2020.
[PDF](#) [arXiv](#)

Summit: Scaling Deep Learning Interpretability by Visualizing Activation and Attribution Summarizations

Fred Hohman, [Haekyu Park](#), Caleb Robinson, and Duen Horng Chau
IEEE Transactions on Visualization and Computer Graphics ([TVCG](#)), Vancouver, BC, Canada, 2020.
[▶ Demo](#) [PDF](#) [arXiv](#)

Visual Analytics for Interpretability on Deep Neural Networks

[Haekyu Park](#), Fred Hohman, Nilaksh Das, Caleb Robinson, and Duen Horng Chau
Women in Machine Learning Workshop ([WiML](#)), co-located with NeurIPS 2019, Vancouver, BC, Canada, 2019.

MLsploit: A Framework for Interactive Experimentation with Adversarial Machine Learning Research

Nilaksh Das, Siwei Li, Chanil Jeon, Jinho Jung, Shang-Tse Chen, Carter Yagemann, Evan Downing, [Haekyu Park](#), Evan Yang, Li Chen, Michael Kounavis, Ravi Sahita, David Durham, Scott Buck, Duen Horng Chau, Taesoo Kim, and Wenke Lee
ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD), [KDD Project](#), Anchorage, Alaska, USA, 2019.
[▶ Demo](#) [PDF](#)

MLsploit: A Cloud-Based Framework for Adversarial Machine Learning Research

Nilaksh Das, Siwei Li, Chanil Jeon, Jinho Jung, Shang-Tse Chen, Carter Yagemann, Evan Downing, [Haekyu Park](#), Evan Yang, Li Chen, Michael Kounavis, Ravi Sahita, David Durham, Scott Buck, Duen Horng Chau, Taesoo Kim, and Wenke Lee
Black Hat Asia - Arsenal, 2019.
[▶ Demo](#) [Abstract](#) [Video](#)

NeuralDivergence: Exploring and Understanding Neural Networks by Comparing Activation Distributions

[Haekyu Park](#) Fred Hohman, and Duen Horng Chau
IEEE Pacific Visualization Symposium ([PacificVis](#)), Bangkok, Thailand, 2019.
[▶ Demo](#) [PDF](#) [arXiv](#)

SIDE: Representation Learning in Signed Directed Networks

Junghwan Kim, [Haekyu Park](#), Ji-Eun Lee, and U Kang
The Web Conference (Previously known as [WWW](#), World Wide Web Conference), Lyon, France, 2018.
[Project](#) [PDF](#)

A Comparative Study of Matrix Factorization and Random Walk with Restart in Recommender Systems

Haekyu Park, Jinhong Jung, and U Kang

IEEE International Conference on Big Data (BigData), Boston, MA, USA, 2017.

Project PDF arXiv

Open-Source Research Projects

CNN Explainer: Learning Convolutional Neural Networks with Interactive Visualization

2020

Keywords: Deep Learning Education, Interactive Visualization, Interactive Animation

Interactive visual system for learning Convolutional Neural Networks.

It was published at IEEE VIS (VAST, TVCG), 2020.

Zijie Jay Wang, Robert Turko, Omar Shaikh, Haekyu Park, Nilaksh Das, Fred Hohman, Minsuk Kahng, Duen Horng (Polo) Chau

► Demo ⚙ Top of Github Trending ★ 4,905 Github stars (as of Oct 2020)

Bluff: Interactively Deciphering Adversarial Attacks on Deep Neural Networks

2020

Keywords: Adversarial Attacks, Neural Network Interpretability, Activation Pathways, Interactive Visual Analytics

Interactive system for visualizing, characterizing, and deciphering adversarial attacks on vision-based neural networks.

It was published at IEEE VIS, 2020.

Nilaksh Das*, Haekyu Park*, Zijie Jay Wang, Fred Hohman, Robert Firstman, Emily Rogers, Duen Horng Chau

(* Equal Contribution)

► Demo

Summit: Scaling Deep Learning Interpretability by Visualizing Activation and Attribution Summarizations

2019

Keywords: Neural Network Interpretability, Attribution Graph, Interactive Visual Analytics

Interactive visualization that scalably summarizes what features a deep learning model has learned and how those features interact to make predictions.

It was published at IEEE VIS (VAST, TVCG), 2019.

Fred Hohman, Haekyu Park, Caleb Robinson, Duen Horng Chau

► Demo

MLsploit: A Framework for Interactive Experimentation with Adversarial Machine Learning Research

2019

Keywords: Adversarial Attacks and Defenses for Machine Learning Models, Interactive Experimentation

User-friendly, cloud-based system that enables researchers and practitioners to rapidly evaluate and compare state-of-the-art adversarial attacks and defenses for machine learning (ML) models.

It was published at a KDD 2019 Project Showcase.

► Demo

SIDE: Representation Learning in Signed Directed Networks

2018

Keywords: Network Embedding, Signed Weighted Directed Graph

General network embedding method that represents both sign and direction of edges in the embedding space.

It was published at the Web Conference (WWW), 2018.

Project

A Comparative Study of Matrix Factorization and Random Walk with Restart in Recommender Systems

2017

Keywords: Recommender System, Matrix Factorization, Random Walk with Restart

We provide a comparative study of matrix factorization and RWR, which are the most representative recommender systems.

It was published at IEEE Big Data, 2017.

Project

Other Projects

DARPA Guaranteeing AI Robustness against Deception (GARD)

2020

Keywords: Defenses for Adversarial Examples, Robustness, Defense using Semantic Coherence

We develop defenses for adversarial attacks on object detector for both RGB images and single-camera video. We augment this object detector to support spatial, temporal, semantic coherence in videos.

RAPIDS and Cybersecurity: A Network Use Case

2019

Keywords: RAPIDS, NVIDIA, GPU-acceleration, Graph, Personalized Page Rank

We showcase an approach to flagging anomalous network communications in a large graph using a combination of structural graph features and graph analytics, running end-to-end in RAPIDS.

Presented at [KDD 2019 NVIDIA RAPIDS tutorial](#) with the [cybersecurity use case notebook](#).

Recommender System for Videos on Oksusu Application

2017

Keywords: Deep Learning, Sequence/Word Embedding, Approx. k-NN, Heterogeneous Features

Our system recommends videos to users of Oksusu application, handling massive data on users' behaviors and heterogeneous information of videos.

SK Telecom, Seoul, Republic of Korea

A Fast Data Compression with Shared Virtual Memory in Heterogeneous System Architecture

2017

Keywords: OpenCL, GPGPU, SVM, HSA

I used general purpose computing on graphics processing units (GPGPU) and Shared Virtual Memory (SVM) in Heterogeneous System Architecture (HSA) for fast data deduplication methods. GPGPU and HSA provide a powerful basis for parallel computing in an easy programmable and efficient way.

Undergraduate thesis

Personalized Recommendation for Credit Card Rewards

2016

Keywords: Coupled Matrix Factorization, Time Series Data

We provide personalized recommendations for credit card rewards to customers using various side information of users and items. The main algorithm is TCMF (Time Coupled Matrix Factorization).

Hyundai Card, Seoul, Republic of Korea

 [News article \(in Korean\)](#)

Talks & Presentation

Bluff: Interactively Deciphering Adversarial Attacks on Deep Neural Networks

Nilaksh Das*, [Haekyu Park](#)*, Zijie Jay Wang, Fred Hohman, Robert Firstman, Emily Rogers, Duen Horng Chau

(* Equal Contribution)

Oct 2020, Oral Presentation, IEEE VIS

Oct 2020, Presentation, Michigan Institute for Data Science (MIDAS) Consortium for researchers in Training

Accelerated Data Science in the Classroom: Teaching Analytics and Machine Learning with RAPIDS

Polo Chau and [Haekyu Park](#)

Mar 2020, Talk, NVIDIA's GPU Technology Conference (GTC)

NeuralDivergence: Exploring and Understanding Neural Networks by Comparing Activation Distributions

Apr 2019, Poster Presentation, PacificVis

Tutorial

RAPIDS and Cybersecurity: A Network Use Case

Keywords: RAPIDS, NVIDIA, GPU-acceleration, Graph, Personalized Page Rank

Presented at KDD 2019 NVIDIA RAPIDS tutorial with the cybersecurity use case notebook

Teaching

Graduate Teaching Assistant

Georgia Institute of Technology, Atlanta, GA
Data and Visual Analytics (CSE 6242)
Instructor: Polo Chau

Fall 2019

Mentoring

Zhiyan Zhou

B.S. in Computer Science, Georgia Institute of Technology
Visualization for Graph Exploration
AI Robustness against Adversarial Attacks

Fall 2020 - Present

Megan Dass

B.S. in Computer Science, Georgia Institute of Technology

AI Robustness against Adversarial Attacks

Fall 2020 - Present

Omar Shaikh

B.S. in Computer Science, Georgia Institute of Technology
Visualization for natural language processing
Received Outstanding Freshman Award

Fall 2019 - Present

Rob Firstman

B.S. in Computer Science, Georgia Institute of Technology
Visualization for deep learning interpretability

Fall 2019 - Spring 2020

Robert Turko

B.S. in Computer Science, Georgia Institute of Technology
Visualization for machine learning education

Fall 2019 - Spring 2020

Licenses and Certifications

NVIDIA DLI Certificate - DLI Platform Course for Instructors, NVIDIA Deep Learning Institute

NVIDIA DLI Certificate - Fundamentals of Deep Learning for Computer Vision, NVIDIA Deep Learning Institute

Professional Service

Reviewer

CIKM 2020
VIS 2020
WiML 2019
KDD 2019
ICML 2019

Professional Membership

The Institute of Electrical and Electronics Engineers (IEEE). Since 2019.

Technical Skills

Programming Languages

Python, JavaScript, HTML, R, Matlab, Java, C, C++, Ocaml, Scheme

Machine Learning / Deep Learning / Data Science

TensorFlow, PyTorch, Keras, scikit-learn, OpenCV, Numpy, Pandas, SciPy, NetworkX

GPU-accelerated Data Science

cuGraph, cuDF, cuML, BlazingSQL, OpenCL

Data Visualization

D3.js, Three.js, WebGL, HoloViews, Matplotlib, WebGL, ggplot